

CMPUT 325

Lecture 2: Fun language

Fun

- List elements are separated by a space
- In general, $f(x_1 \dots x_n)$ --> x , notice the spacing for parameters
- `first(L)` returns first element of L , error if L is not a list or an empty list
- `rest(L)` returns L without the first element in L , error if L is not a list or an empty list
- Using compositions of calls to `first` (or `f`) and `rest` (or `r`), we can get any component (atom or sublist) from a given list, regardless of its depth
- For brevity, $f(f(r(L)))$ can be written as `ffr(L)`
- `cons(x, L)` returns a new list $K = (x \ y_1 \ y_2 \ y_3 \dots y_n)$ where x is an atom and $L = (y_1 \ y_2 \ y_3 \dots y_n)$
- In theory, there is a way to define any primitive function from scratch.
- Assume we have the following primitive functions:
 - Arithmetic: `+`, `-`, `*`, `/`, and so on
 - Comparators
 - `if`, `then`, `else`
 - `null(x)`
 - `true` iff x is an empty list
 - `atom(x)`
 - `true` iff x is an atom
 - `eq(x, y)`
 - `true` iff x and y are the same atom
 - `eq(a, a)` returns `true`
 - `eq(a, b)` returns `false`
 - `eq((a), (a))` returns `false` because the arguments are lists
- Similar idea is used to create RISC syntax.
- Remarks:
 - No notion of variable-as-storage
 - No assignment statement as in procedural languages
 - No loop constructs
 - **Recursion** as the only mechanism to define non-trivial functions
- 1-step evaluation = replacement + substitution
- Examples covered:
 - `length(L)`, `append(L1, L2)`, `last(L)`, `removeLast(L)`
- It makes sense to break the solution into smaller functions
- One function must do only one thing

Lecture 2: Intro to LISP

- `reverse(L)` using `append(L1, L2)`
- Abstract data type for binary tree
 - Goal: implement a binary tree and some operations, such as inserting elements
 - Two main tasks:
 - Decide how trees are represented by lists
 - Implement an abstract data type for binary trees and the operations on them
 - User will work with trees using only these functions. The user is protected from the details of our data representation
 - Bottom up construction
- Binary tree representation
 - Empty tree: `nil`
 - Nonempty tree: `(left-subtree, node-value, right-subtree)`
 - Selectors:
 - `leftTree(Tr) = f(Tr)`
 - `rightTree(Tr) = f(r(r(Tr)))`
 - `nodeValue(Tr) = f(r(Tr))`
 - Constructors:
 - `consNilTr() = nil`
 - return an empty tree
 - `consTree(L, V, R) = cons(L, cons(V, cons(R, nil)))`
 - construct tree with subtrees L, R and value V
 - Test:
 - `isEmpty(Tr) = eq(Tr, nil)`
 - return true iff Tr is an empty tree
- Building an abstract data type
 - Functions are the only ones that need direct knowledge of our tree representation
 - Everything else can be implemented in terms of these basic functions - providing such a base set of functions is the essence of implementing an abstract data type in functional programming
- `insert` into tree
 - assume our trees contain integer values and are sorted such that every value in the left subtree $<$ node value $<$ all values in right subtree
 - unique values
 - `insert(Tr, int)` inserts `int` into the binary tree `Tr`

```

insert(Tr, int) =
  if isEmpty(Tr)
    then consTree(consNilTr(), Int, consNilTr())
  else if Int = nodeValue(Tr)
    then Tr
  else if Int < nodeValue(Tr)
    then consTree(insert(leftTree(Tr), Int, nodeValue, Tr),
rightTree(Tr))
  else consTree(leftTree(Tr), nodeValue(Tr), insert(rightTree(Tr),
Int))

```

LISP

- interpreted language
- case **insensitive**
- uses read-eval-print-loop (REPL) similar to a shell such as bash
 - read input
 - evaluate input
 - print result of evaluation
 - loop back to beginning
- Functions are defined by

```
(defun function-name parameter-list body)
```

- Example:
 - Function definition: (defun plus (x y) (+ x y))
 - Function application: (plus 3 4)
- Lisp always interprets (e1 e2 e3 ...) as a function application. Use quote to "atomify" the expression
- An empty list is represented by either () or nil. Both are considered the same atom in Lisp.
- nil also represents false
- T represents true

Lecture 4

- (if condition then-part else-part) is a special function because not every block is run, unlike other functions
- trace to see calls and returns to specific functions
- untrace stops the tracing
- functions can take variable number of arguments
- (let ((x 3) (y 4)) (* (+ x y) x)) evaluates expression, but replaces names x and y with their values 3 and 4
- let does not allow using one variable to define another, use let* instead
- eq is true iff both are equal *atoms*, runs in a single machine instruction

- equal is more general

```
(cond (P1 S1)
      (P2 S2)
      (P3 S3)
      ...
      (T Sn)
)
```

- General form of cond (do not use it):

```
(cond (P1 S11 S12 ... S1m)
      (P2 S21 S22 ... S2m)
      ...
      (T Sn1 Sn2 ... Snm)
)
```

- If P1 is true then evaluate S11, S12, ... and S1m and return the result of evaluating S1m
- list
- caar, cddadr, etc.
- print and format for printing, strings
- random
- Use quote when everything is constant
- Use list when some contents are the result of evaluating functions
- Use cons
 - for the result in recursion when you have computed a first element and the rest of a list
 - for dotted pairs
- (car (cdr (car (cdr (cdr (cdr L)))))) = (cadr (caddr L))
 - max 4 levels deep
- Simple printing (print arg)
- Formatted printing: (format t format-string arg1 ...)
- (random N) generates a uniformly random integer from 0..N-1 if N is an integer
- (random F) generates a uniformly random floating point number in range [0..F)
- Accumulators
 - helper function with an extra parameter
 - the extra parameter accumulates the required result
 - Issues with simple recursion:
 - no real computation until hits the base case

- all computation happens on return from recursion
- Example 1: reverse using an accumulator

```
(defun reverse_helper (L ResultSoFar)
  (if (null L)
      ResultSoFar
      (reverse_helper (cdr L)
                      (cons (car L) ResultSoFar)))
  )

(defun reverseAC (L)
  (reverse_helper L nil)
  )
```

- Comparing accumulator with standard recursion
 - Standard recursion on a list
 - recurse to the end of list
 - compute result on return from recursion
 - bottom-up computation
 - Accumulators
 - accumulates results-so-far
 - computes results top-down
 - needs an extra accumulator variable for partial result
 - Questions to think about to decide whether accumulators should be used
 - top-down or bottom-up?
- Programming loops in LISP
 - In pure functional programming, we use recursion instead of loops although LISP has loop constructs (for, do, loop, ...)
 - break loop into two steps:
 - what to do in each run through the loop
 - how to solve the rest of the problem by recursion

Lecture 5

- Symbolic expressions (S-Expressions, s-expr, sexpr)
 - universal data structure for Lisp
 - generalization of atoms and lists
 - all atoms and lists are sexpr
 - but not all atom and lists are sexpr
 - dotted pair: (x . y)

- Definition

- atom is an s-expression
- if x_1, x_2, \dots, x_n are s-expressions then $(x_1 \dots x_n)$ is an s-expression
- if x_1 and x_2 are s-expressions, then $(x_1 . x_2)$ is an s-expression (a dotted pair)
- Examples:
 - hello
 - (a b c)
 - (a (b) (()))
 - (a . b)
 - (a . (b . c))
 - (1 2 3 (4 . 5))

- `(car (x . y))` returns x

- `(cdr (x . y))` returns y


- `(car (cons 'x 'y)) = x`

- `(cdr (cons 'x 'y)) = y`


- `.` must be surrounded by whitespace: `(a.b)` is a list containing the atom `a.b`

- machine-level representations

- Example 1: `(cons 1 2)` or `(1 . 2)`




- Example 2: ((1 . (2 . 3)) . 4)



```
(cons (cons 1
            (cons 2 3))
      4)
```

- Example 3: List representation (1 2 3 4)



- (a . nil) = (a) because (cdr '(a . nil)) = (cdr ('(a))) = nil
- Every list can be written as nested dotted pairs:
(1 . (2 . (3 . (4 . nil))))
- Why use dotted pairs?
 - saves memory
 - simplifies direct access

Lecture 6

- Higher order functions
 - Definition: a function that takes other function(s) as input and/or produce function(s) as output
 - often used to separate:
 - a computation pattern
 - specific repeated action

- Example 1:
 - Pattern: iterate over a list
 - Action: Compute the same function for each list element
- Example 2:
 - Pattern: reduce a list to a single result
 - Action: reduce two arguments to one
- Some typical higher order functions
 - Map - apply some function to all elements of a list
 - Reduce - apply two argument function repeatedly
 - Filter - select list elements that pass a test
 - Vector - apply many functions to one element
- mapcar in Lisp

```
(defun plus1 (x) (+ x 1))
(mapcar 'plus1 '(1 2 3 4 5)) --> (2 3 4 5 6)
```

- reduce - general definition
 - input:
 - function g
 - function's identity id
 - a list L = (a1 a2 ... an)
 - compute: (g a1 (g a2 ... (g an id) ...))
 - example in Lisp:
 - (reduce '* '(2 6 4)) --> 48
- Why define high order functions?
 - use case: a common computation pattern, where the details can vary
 - removing code duplication
 - apply and funcall tells Lisp that there is a function to be called
 - only differs in syntax, same functionality
 - (apply function-name (arg1 ... argn))
 - (funcall function-name arg1 ... argn)

Lecture 7

Lambda Functions

- get rid of named functions, why?
 - a function was the result of a higher order function
 - tried to return this newly computed function
- lambda functions are **function definitions** without names

- Syntax: (lambda (x1 ... xn) body)
- Example: ((lambda (x y) (+ x y)) 5 3)
- lambda function application

- to apply a lambda function

```
((lambda (x1 ... xn) body) a1 ... an)
x1 ... xn are formal arguments of function
a1 ... an are actual parameters for which we want to
evaluate the function
```

- Lisp-1 and Lisp-2
 - Lisp-1 systems: "values" and functions in the same namespace
 - Lisp-2 systems: in separate spaces
 - Common Lisp standard requires Lisp-2
 - if we have a variable bound to a function, we need to tell SBCL this is a function to be called
 - Consequence:
 - Working with lambda functions is much messier in Lisp-2 systems than in Lisp-1
 - never quote a lambda expression in Lisp-2
- function
 - syntax: (function arg)
 - purpose: evaluates lambda function given by arg
 - takes lambda function as its argument
 - returns function definition in an internal format used by SBCL
 - compiles it and returns an internal representation of the compiled code
 - representation is called **closure**
 - next, the function in the closure can be called in an application
 - use funcall or apply for the application
- function vs funcall vs apply
 - function takes as argument a *function definition* and returns an internal representation of that definition
 - does NOT apply the function
 - funcall and apply are for *function application*
- applying lambda functions
 - use funcall and apply as usual by giving the whole lambda function as an argument

Lambda Calculus

- Intro to lambda calculus
 - formal, abstract language
 - all functions are defined without giving them names
 - Lisp is based on lambda calculus but adds a large language on top of it

- formal language with only four concepts:

```

[identifier] := a | b | ...
[function]   := (lambda (x) [expression])
[application] := ([expression] [expression])
[expression] := [identifier] | [application] | [function]

```

- identifier: corresponds to an atom
 - function: is a lambda function definition
 - expression: can be an arbitrary lambda expression. It plays the role of the body in the function definition.
 - corresponds to sexpr in Lisp
 - application: both function and arguments can be any expressions
 - usually, first expression will evaluate to a lambda function and the second argument will evaluate to the arguments for the lambda function
- all valid expressions defined by this language are called lambda expressions
- lambda expressions can represent any computation
- Unary vs N-ary functions
 - [function] := (lambda (x) [expression])
 - we only have unary functions - functions that take **one** parameter
 - any **n-ary function** (function with n arguments) can be defined using **a series of unary functions**
 - consequence:
 - to understand the model of computation for general functional programming
- Curried functions
 - Goal: define an n-ary functions by a series of unary functions
 - can solve this by using higher order functions
 - main idea:
 - series of n unary function applications
 - each application processes one argument
 - the application produces a new function which has this argument hardcoded
 - Intuition
 - Example: (plus 5 2) is a function with two args
 - (plus5 2) is a function with one argument, the "add 5" is hardcoded into the new function plus5
 - function takes only the first argument, and produces as result a new function
 - this new function now takes the second argument and produces a new function
 - the function that takes the last argument will have other argument

- values "hardcoded"
 - each function is computed on the fly by all the previous function applications
- splitting up the two argument function

```
(lambda (x y) (if (< x y) x y))
```

is equivalent to

```
(lambda (x) (lambda (y) (if (< x y) x y)))
```

Lecture 8

Reductions in Lambda Calculus

- Goal: reduce a lambda expression to its **simplest possible form**
- This process is called *operational semantics* of lambda calculus
- In lambda calculus, computation is **the process of reductions from one expression to another expression**
- Example:

```
((lambda (x) (x 2)) (lambda (z) (+ z 1))) → (+ 2 1)
```

- Shorthand notation in lambda calculus
($\lambda x. (+ x 1)$) for

```
(lambda (x) (+ x 1))
```

- In lambda calculus, we do not need any of the primitive functions
 - numbers can be represented by lambda expressions
- Beta reduction
 - function application
 - written as \rightarrow^{β}
 - rule:
 - given an expression $((\lambda x) \text{ body }) a$, reduce it to body
 - replace all occurrences of x in body by a
 - Example:
 - $((\lambda x) (x (x 1))) 5 \rightarrow^{\beta} (5 (5 1))$
 - $(\lambda x) (x x x x x) (1 2 3 4 5) \rightarrow^{\beta} ((1 2 3 4 5) (1 2 3 4 5) (1 2 3 4 5) (1 2 3 4 5) (1 2 3 4 5))$
 $x = (1 2 3 4 5)$
- Alpha reduction


- \rightarrow^{α} means renaming variables
- Intuition: changing the name of local variables in a function does not change the meaning
- *name conflict* between arguments:
 - `(defun f(x x) (- x x))`
 - this gives a compile-time error in SBCL: variable x occurs more than once in a lambda expression
- Example: `(lambda (x) (+ x y))`
 - x is **bound** in the scope of `(lambda (x) ...)`
 - y is **free**
 - x can be renamed to anything except y
 - y cannot be renamed
- Free vs bound variables
 - depend on their scopes
 - like global and local variables
- In lambda calculus, a bound variable's name can be replaced by another if the latter does not cause any name conflict
 - use a new variable name
 - called alpha reduction
 - without alpha reduction, direct substitution does not always work
- **Note:** Perform alpha reduction first!
 - $((\lambda x (\lambda z (x z))) z)$
 - rename the zz in $(\lambda z \dots)$
 - $((\lambda x (\lambda u (x u))) z)$
 - now the bound variable is called u and will not conflict with the argument z
 - finally replace x by z in body
 - $(\lambda (u) (z u))$
- Scope of variables and beta-reduction
 - scope of a variable should be preserved by variable renaming to ensure that reduction is correct
 - $((\lambda x (\lambda z (x z))) z) \rightarrow^{\beta} (\lambda u (z u))$
 - where u is some new variable
 - exercise: fill in the steps
 - correct beta reductions can always be achieved by renaming (alpha-reduction), if needed
 - beta-reduction using direct substitution
- Summary of reductions

- one β -reduction corresponds to a one-step function application
- direct substitution does not always work, variables may need to be renamed before the substitution step
- the substitution of the formal variable by the argument must be done carefully to avoid name conflicts
- α -reduction renames function arguments
- after using such renaming where necessary, a simple substitution in the body gives a correct beta-reduction
- to be safe can always use α -reduction with names for bound variables

Lecture 9

Normal form, order of reduction and Church Rosser Theorem

- **Normal form:** a lambda expression that cannot be reduced further (by beta reduction)
- not all lambda expressions have a normal form
 - apply a sequence of reductions and the same expression will be obtained
- $((\lambda x. (x x)) (\lambda z. (z z)))$ does not have a normal form
- order of reduction
 - Normal order reduction (NOR):
 - evaluate leftmost outermost application
 - Applicative order reduction (AOR):
 - evaluate leftmost innermost application
- tie-breaking rules:
 - choose the leftmost one for either order when dealing with multiple choices
- Efficiency
 - AOR is generally more efficient
 - NOR terminates more often, AOR might get into infinite reduction
- Church Rosser theorem
 - Church and Rosser proved two important properties of reductions and normal forms
 - If A reduces to B and A reduces to C then there exists an expression D such that B reduces to D and C reduces to D
 - If A has a normal form E then there is a **NOR** from A to E
 - Remarks
 - If A reduces to B and A reduces to C then there exists an expression D such that B reduces to D and C reduces to D



- No matter what reduction strategies are used initially to get from A to B and C, there is always a way to converge from both B and C back to the same expression D
- There is always at most one normal form
- If A has a normal form E then there is an NOR from A to E
- NOR guarantees termination if the given expression has a normal form
 - can be inefficient
- Theorem does not tell us if a normal exists

Lecture 10: Interpreter based on context and closure

- implement an interpreter for a Lisp-like language based program, only consisting of lambda functions
- interpreter based on two concepts: **context** and **closure**
- Language
 - variables
 - constant expressions: (quote e)
 - arithmetic
 - relations and logic
 - primitives for s-expr:
 - car, cdr, cons, atom, null
 - if
 - lambda functions
 - function call
 - simple block (let (x1.e1) ... (xk.ek) e)
 - (optional) recursive block
 - (letrec (x1 . e1) ... (xk . ek) e), instead of the Lisp way to simplify
- Why not just use β -reductions?

- determine the scope of each parameter
 - detect potential name conflicts
 - implement variable renaming (α -reductions)
 - implement direct substitutions
 - it is possible to do it, but not very efficient
 - main problem:
 - need to check all the above repeatedly after each substitution step
- new idea
 - key idea: *delay* the substitutions by using contexts and closures
 - this technique is used in real Lisp interpreters
 - also helps with understanding compilation
- **Context**
 - a context is *a list of bindings*

$$[n_1 \rightarrow v_1 \dots n_k \rightarrow v_k]$$
 - where n_i are identifiers, v_i are expressions (value that the variable is bound to)
 - v_i can also be a closure, represents the state of an incomplete evaluation
 - a context is used to keep record of lookup name bindings
 - a context can be *extended* when a new binding $n \rightarrow v$ is created in a function application
- **evaluation with a context**
 - always begin with an empty context
 - compare with other programming languages
 - may have global variables already bound to values
 - in the middle of evaluating an expression, the context is usually non-empty
 - substitutions are delayed to the point where the value of a variable is really needed for the evaluation to continue
 - variables are left as-is
 - a variable is bound as "needed"...
 - if binding can be found in context
- **Closure**
 - pair $[f, CT]$
 - f is a lambda function
 - CT is a (possibly empty) context
 - when function f is applied we know the parameters and body
 - we get values for the variables **from the context** which are used in the body of f
- **Function application in a context - Algorithm**
 - when interpretation of a program starts, the context is empty

- when a function is applied:
 - evaluate the arguments in the current context
 - evaluate the functional part in the current context
 - extend the context
 - bind the parameter names to the evaluated arguments
 - add these bindings to current context to form the next context
 - evaluate the body of the function in this extended context
- **Implementation of context for interpreter**
 - define a data structure to represent a context
 - two lists, name list and value list
 - both lists are in sync
 - for each name there is a corresponding value stored in the same location in the other list
 - name list is a list of lists of atoms
 - value list is a list of lists of symbolic expressions
- example
 - name list `((x y) (z) (w s))`
 - value list `((1 2) ((lambda (x) (* x x)))`
 - list of three sublists
 - corresponds to three (nested) lambda function applications
 - in previous notation, this implements the context
- compare context and closure model with runtime execution model of programs
 - compare to runtime model of a programming language
 - a runtime model has:
 - a call stack for all active functions
 - one stack frame for each function
 - the names are compiled away, replaced by relative addresses on the stack
 - in our model, we still keep the names in the context
 - we will soon see a similar compilation for Lisp, also compiling away names, in the SECD machine
- **name lookup for context**
 - search for a name in a context
 - walk synchronously over both name and value lists
 - if a name is found:
 - s-expr in the same position in the value list is its binding
 - name lookup function `assoc(x, n, v)`
 - name list n

- value list v
- name to lookup stored in x
- linear search implementation

Lecture SECD machine

- programming language implementations:
 - interpreter
 - compiler to assembly, real machine code
 - compiler to virtual machine
- SECD is a virtual machine that runs simple Lisp programs
- virtual machines
 - not real hardware, but has its own set of instructions
 - to run, implement on a real machine
- SECD usage
 - compiler from Lisp to SECD machine code
 - execute the compiled code on an abstract cmachine
 - main differences to interpreter
 - once compiled, code can be executed many times
 - code optimization is possible
- SECD Machine
 - consists of four stacks
 - s: Stack used for evaluation of expressions
 - e: Environment stores the current value list
 - c: Control stores the machine instructions
 - d: Dump stores *suspended* invocation context
- SECD operations and functions
 - NIL: push a nil pointer
 - LD: load from environment
 - LCD: load constant
 - LDF: load function
 - AP: apply function
 - RTN: return
 - SEL: select in if statement
 - JOIN: rejoin main control
 - builtin funtions: +, *, ATOM, CAR, CONS, EQ, etc.
- definition of SECD operations
 - each operation is defined by its effect on the four stacks
 - s e c d ---> s' e' c' d'
 - representation of a stack s

- s-expression with dot notation
- note: for brevity, in this topic spaces around the '.' are omitted as required in Lisp
- first position (car s) = top of the stack
- push onto stack s ---> (e.s)
- pop from the stack (e.s) ---> s